

House Price Prediction

November 24, 2019

1 Problem Statement

Ask a home buyer to describe their dream house, and they probably won't begin with the height of the basement ceiling or the proximity to an east-west railroad. But this dataset proves that much more influences price negotiations than the number of bedrooms or a white-picket fence.

- You can use sklearn, numpy pandas, matplotlib
- Apply linear regression, random forest regressor
- Steps include
 - Read the data using pandas dataframe
 - Do feature engineering and separate target value price
 - Remove those features which are not contributing
 - Apply train test split on given data
 - Apply ML algorithm using sklearn on training data and do prediction of test set
 - Evaluation metric should be RMSE

Follow the Data Science pipeline and build a Machine Learning model that will predict House Prices based on given features. The pipeline is as follows:

1. Data Wrangling and Preprocessing
2. Exploratory Data Analysis
3. Feature Selection
4. Model Training
5. Testing and Optimization

Bonus part: Do detail Exploratory data analysis to get good features.

```
In [17]: # If you want to install any missing packages, then uncomment the lines given below a
        # to ensure that you have all the dependencies you need to run the notebook.
        #import sys
        #{sys.executable} -m pip install xgboost
        from scipy import stats

In [2]: # Libraries
import pandas as pd
import numpy as np
import missingno as mno
import seaborn as sns
import matplotlib.pyplot as plt
```

1.1 Data Description

The first thing you need to do before solving any Data Science problem is getting familiar with the dataset. Get to know your data by printing out some stats, checking its dimensions and checking data types of features.

```
In [4]: # Load training data
```

```
data = pd.read_csv('house_data.csv')
data.head()
```

```
Out[4]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	\
0	7129300520	20141013T000000	221900.0	3	1.00	1180	
1	6414100192	20141209T000000	538000.0	3	2.25	2570	
2	5631500400	20150225T000000	180000.0	2	1.00	770	
3	2487200875	20141209T000000	604000.0	4	3.00	1960	
4	1954400510	20150218T000000	510000.0	3	2.00	1680	

	sqft_lot	floors	waterfront	view	...	grade	sqft_above	\
0	5650	1.0	0	0	...	7	1180	
1	7242	2.0	0	0	...	7	2170	
2	10000	1.0	0	0	...	6	770	
3	5000	1.0	0	0	...	7	1050	
4	8080	1.0	0	0	...	8	1680	

	sqft_basement	yr_built	yr_renovated	zipcode	lat	long	\
0	0	1955	0	98178	47.5112	-122.257	
1	400	1951	1991	98125	47.7210	-122.319	
2	0	1933	0	98028	47.7379	-122.233	
3	910	1965	0	98136	47.5208	-122.393	
4	0	1987	0	98074	47.6168	-122.045	

	sqft_living15	sqft_lot15
0	1340	5650
1	1690	7639
2	2720	8062
3	1360	5000
4	1800	7503

```
[5 rows x 21 columns]
```

```
In [5]: # Dimensions of training data
```

```
data.shape
```

```
Out[5]: (21613, 21)
```

```
In [6]: # Explore columns
```

```
data.columns
```

```
Out[6]: Index(['id', 'date', 'price', 'bedrooms', 'bathrooms', 'sqft_living',  
              'sqft_lot', 'floors', 'waterfront', 'view', 'condition', 'grade',
```

```
'sqft_above', 'sqft_basement', 'yr_built', 'yr_renovated', 'zipcode',
'lat', 'long', 'sqft_living15', 'sqft_lot15'],
dtype='object')
```

```
In [7]: # Description
data.describe()
```

```
Out[7]:
```

	id	price	bedrooms	bathrooms	sqft_living \
count	2.161300e+04	2.161300e+04	21613.000000	21613.000000	21613.000000
mean	4.580302e+09	5.400881e+05	3.370842	2.114757	2079.899736
std	2.876566e+09	3.671272e+05	0.930062	0.770163	918.440897
min	1.000102e+06	7.500000e+04	0.000000	0.000000	290.000000
25%	2.123049e+09	3.219500e+05	3.000000	1.750000	1427.000000
50%	3.904930e+09	4.500000e+05	3.000000	2.250000	1910.000000
75%	7.308900e+09	6.450000e+05	4.000000	2.500000	2550.000000
max	9.900000e+09	7.700000e+06	33.000000	8.000000	13540.000000

	sqft_lot	floors	waterfront	view	condition \
count	2.161300e+04	21613.000000	21613.000000	21613.000000	21613.000000
mean	1.510697e+04	1.494309	0.007542	0.234303	3.409430
std	4.142051e+04	0.539989	0.086517	0.766318	0.650743
min	5.200000e+02	1.000000	0.000000	0.000000	1.000000
25%	5.040000e+03	1.000000	0.000000	0.000000	3.000000
50%	7.618000e+03	1.500000	0.000000	0.000000	3.000000
75%	1.068800e+04	2.000000	0.000000	0.000000	4.000000
max	1.651359e+06	3.500000	1.000000	4.000000	5.000000

	grade	sqft_above	sqft_basement	yr_built	yr_renovated \
count	21613.000000	21613.000000	21613.000000	21613.000000	21613.000000
mean	7.656873	1788.390691	291.509045	1971.005136	84.402258
std	1.175459	828.090978	442.575043	29.373411	401.679240
min	1.000000	290.000000	0.000000	1900.000000	0.000000
25%	7.000000	1190.000000	0.000000	1951.000000	0.000000
50%	7.000000	1560.000000	0.000000	1975.000000	0.000000
75%	8.000000	2210.000000	560.000000	1997.000000	0.000000
max	13.000000	9410.000000	4820.000000	2015.000000	2015.000000

	zipcode	lat	long	sqft_living15	sqft_lot15
count	21613.000000	21613.000000	21613.000000	21613.000000	21613.000000
mean	98077.939805	47.560053	-122.213896	1986.552492	12768.455652
std	53.505026	0.138564	0.140828	685.391304	27304.179631
min	98001.000000	47.155900	-122.519000	399.000000	651.000000
25%	98033.000000	47.471000	-122.328000	1490.000000	5100.000000
50%	98065.000000	47.571800	-122.230000	1840.000000	7620.000000
75%	98118.000000	47.678000	-122.125000	2360.000000	10083.000000
max	98199.000000	47.777600	-121.315000	6210.000000	871200.000000

```
In [8]: # Check Datatypes
data.dtypes
```

```

Out[8]: id                int64
        date              object
        price             float64
        bedrooms          int64
        bathrooms         float64
        sqft_living       int64
        sqft_lot          int64
        floors            float64
        waterfront        int64
        view              int64
        condition         int64
        grade             int64
        sqft_above        int64
        sqft_basement     int64
        yr_built          int64
        yr_renovated      int64
        zipcode           int64
        lat               float64
        long              float64
        sqft_living15     int64
        sqft_lot15        int64
dtype: object

```

1.2 Data Wrangling and Preprocessing

We have to preprocess our data in order to make it useful for data analysis and model training. Although, the steps involved vary depending on the problem and the dataset but here we have provided a roughly generic approach which is applicable for most problems. The steps involved are as follows:

1. Look for Null or Missing Values
2. Change data type of features, if required
3. Encode data of categorical features
4. Deal with Null or Missing values

```

In [9]: # Check for any null or missing values
        data.isnull().values.any()

```

```

Out[9]: False

```

```

In [10]: # Check missing values in each column of training data
         data.isnull().sum()

```

```

Out[10]: id                0
         date              0
         price             0
         bedrooms          0
         bathrooms         0
         sqft_living       0

```

```

sqft_lot      0
floors        0
waterfront    0
view          0
condition     0
grade         0
sqft_above    0
sqft_basement 0
yr_built      0
yr_renovated  0
zipcode       0
lat           0
long          0
sqft_living15 0
sqft_lot15    0
dtype: int64

```

Thus we can see that there are no missing values in the data.

```

In [14]: # Convert object type to date type
data['date'] = pd.to_datetime(data['date'])
data.head()

```

```

Out[14]:
   id      date      price  bedrooms  bathrooms  sqft_living  \
0  7129300520 2014-10-13 221900.0         3         1.00         1180
1  6414100192 2014-12-09 538000.0         3         2.25         2570
2  5631500400 2015-02-25 180000.0         2         1.00          770
3  2487200875 2014-12-09 604000.0         4         3.00         1960
4  1954400510 2015-02-18 510000.0         3         2.00         1680

   sqft_lot  floors  waterfront  view  ...  grade  sqft_above  \
0      5650     1.0           0     0  ...     7         1180
1      7242     2.0           0     0  ...     7         2170
2     10000     1.0           0     0  ...     6          770
3       5000     1.0           0     0  ...     7        1050
4       8080     1.0           0     0  ...     8        1680

   sqft_basement  yr_built  yr_renovated  zipcode    lat    long  \
0              0     1955              0    98178  47.5112 -122.257
1             400     1951            1991    98125  47.7210 -122.319
2              0     1933              0    98028  47.7379 -122.233
3             910     1965              0    98136  47.5208 -122.393
4              0     1987              0    98074  47.6168 -122.045

   sqft_living15  sqft_lot15
0             1340         5650
1             1690         7639
2             2720         8062

```

3	1360	5000
4	1800	7503

[5 rows x 21 columns]

1.3 Data Analysis and Visualizations

Performing a detailed analysis of the data helps you understand which features are important, what's their correlation with each other which features would contribute in predicting the target variable. Different types of visualizations and plots can help you achieve that. These include:

1. Bar Plots
2. Joint Plots
3. Box Plots
4. Correlation Heatmap
5. Distribution Plot
6. PCA Bi-plot

The 'Target Variable' in this data is the price column. We will now perform some analysis on the target variable to get a better insight into what we are working on.

```
In [18]: plt.subplots(figsize=(12,9))
         sns.distplot(data['price'], fit=stats.norm)

         # Get the fitted parameters used by the function



         (mu, sigma) = stats.norm.fit(data['price'])

         # plot with the distribution

         plt.legend(['Normal dist. ($\mu=${:.2f}$ and $\sigma=${:.2f}$ )'.format(mu, sigma)], 1)
         plt.ylabel('Frequency')

         #Probability plot

         fig = plt.figure()
         stats.probplot(data['price'], plot=plt)
         plt.show()
```



This target variable is right skewed. Now, we need to transform this variable and make it normal distribution.

Here we use log for target variable to make more normal distribution


```
In [19]: #we use log function which is in numpy
data['price'] = np.log1p(data['price'])


#Check again for more normal distribution
plt.subplots(figsize=(12,9))
sns.distplot(data['price'], fit=stats.norm)

# Get the fitted parameters used by the function
(mu, sigma) = stats.norm.fit(data['price'])

# plot with the distribution
plt.legend(['Normal dist. ( $\mu$ = $ {:.2f} and  $\sigma$ = $ {:.2f} )'.format(mu, sigma)], 1)
plt.ylabel('Frequency')

#Probability plot
fig = plt.figure()
stats.probplot(data['price'], plot=plt)
plt.show()
```





Correlation between train attributes

```
In [20]: # Separate variable into new dataframe from original dataframe which has only numeric
# There are 20 variables of numerical type
train_corr = data.select_dtypes(include=[np.number])
```

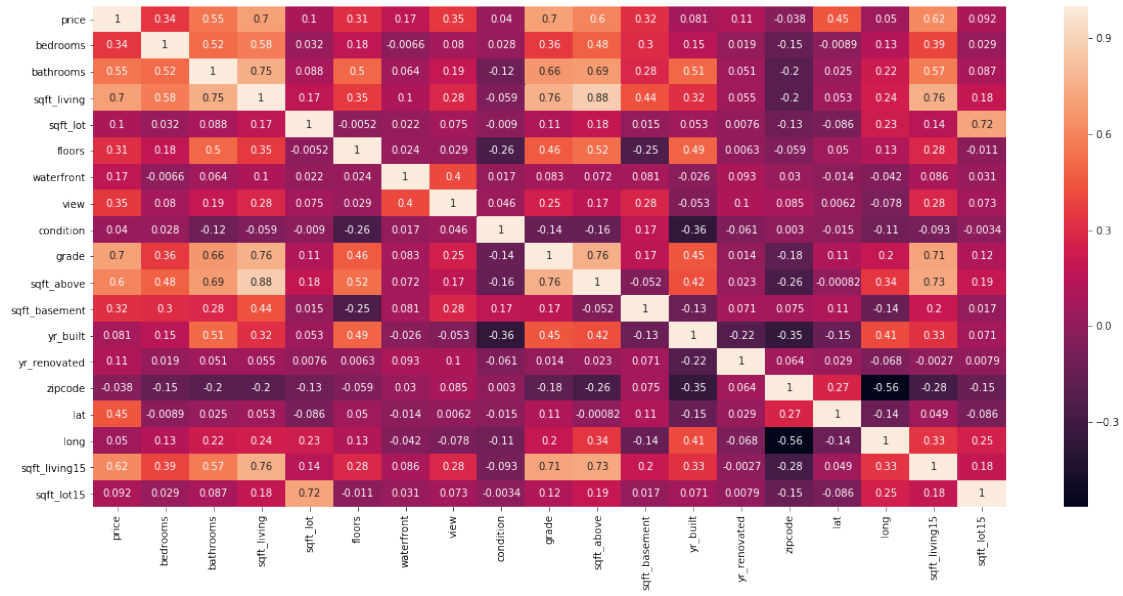
```
In [21]: train_corr.shape
```

```
Out[21]: (21613, 20)
```

```
In [22]: # Delete Id because that is not need for correlation plot
del train_corr['id']
```


```
In [23]: # Correlation plot
corr = train_corr.corr()
plt.subplots(figsize=(20,9))
sns.heatmap(corr, annot=True)
```

```
Out[23]: <matplotlib.axes._subplots.AxesSubplot at 0x1c85d41e4a8>
```



Top 50% Correlation train attributes with price

```
In [24]: top_feature = corr.index[abs(corr['price']>0.5)]
plt.subplots(figsize=(12, 8))
top_corr = data[top_feature].corr()
sns.heatmap(top_corr, annot=True)
plt.show()
```



```
In [34]: print("Find most important features relative to target")
         corr = data.corr()
         corr.sort_values(['price'], ascending=False, inplace=True)
         corr.price
```

Find most important features relative to target

```
Out[34]: price          1.000000
         grade          0.703634
         sqft_living    0.695341
         sqft_living15  0.619312
         sqft_above     0.601802
         bathrooms     0.550802
         lat            0.449174
         view           0.346522
         bedrooms       0.343561
         sqft_basement  0.316970
         floors         0.310558
         waterfront     0.174586
         yr_renovated    0.114498
         sqft_lot       0.099622
```

```
sqft_lot15      0.091592
yr_built        0.080654
long            0.049942
condition       0.039558
id             -0.003819
zipcode        -0.038306
Name: price, dtype: float64
```


Here grade is highly correlated with target feature of price by 71%

```
In [30]: #unique value of sqft_living
data.grade.unique()
```

```
Out[30]: array([ 7,  6,  8, 11,  9,  5, 10, 12,  4,  3, 13,  1], dtype=int64)
```


```
In [31]: sns.barplot(data.grade, data.price)
```

```
Out[31]: <matplotlib.axes._subplots.AxesSubplot at 0x1c86d6eee80>
```



```
In [32]: #boxplot
plt.figure(figsize=(18, 8))
sns.boxplot(x=data.grade, y=data.price)
```

```
Out[32]: <matplotlib.axes._subplots.AxesSubplot at 0x1c86d615b00>
```

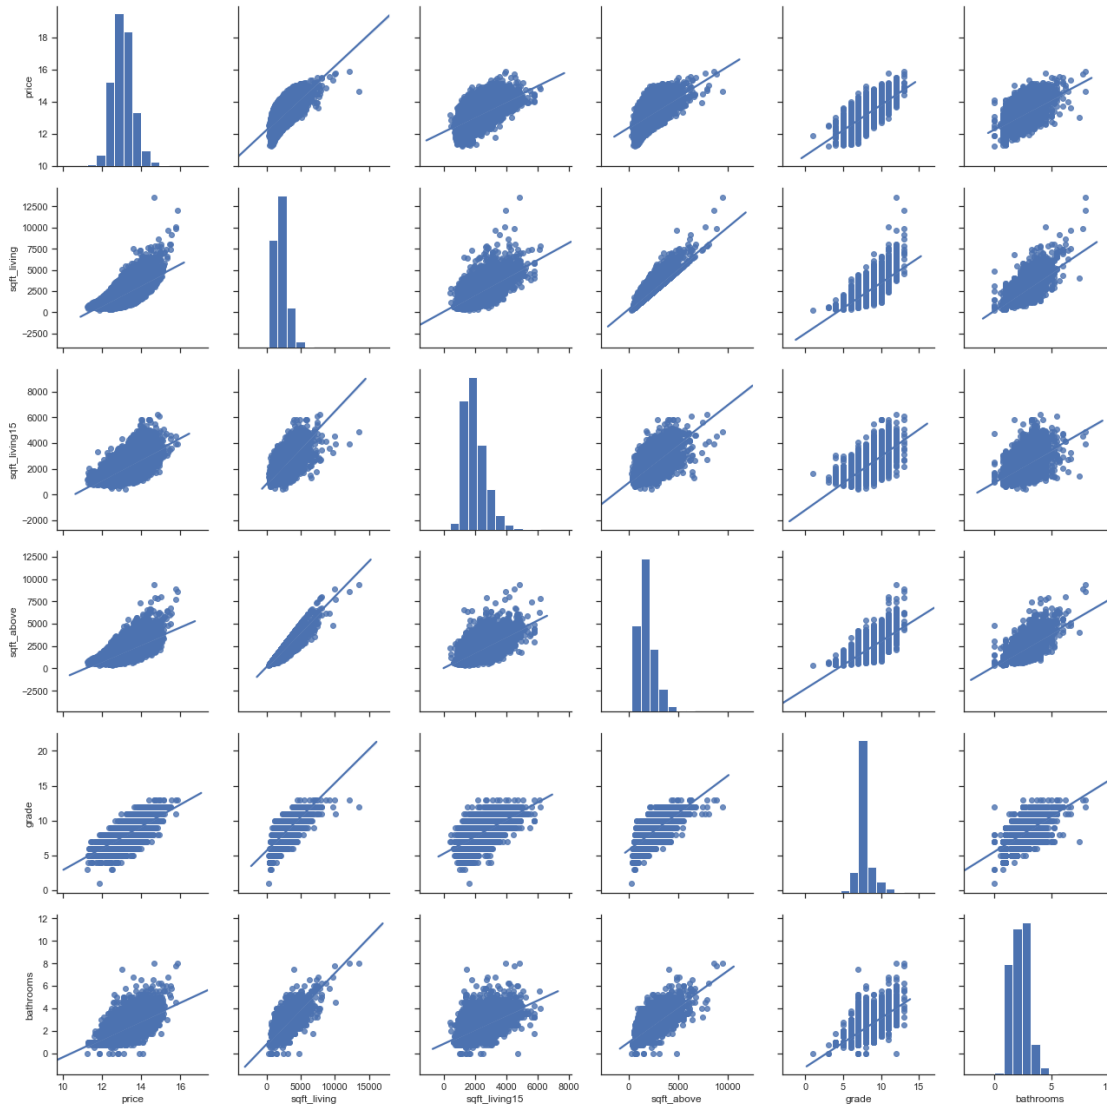


Evaluation of the top_features which are contributing to the price of the house.

```
In [28]: col = ['price', 'sqft_living', 'sqft_living15', 'sqft_above', 'grade', 'bathrooms']
sns.set(style='ticks')
sns.pairplot(data[col], size=3, kind='reg')
```

C:\Users\alina\Anaconda3\lib\site-packages\seaborn\axisgrid.py:2065: UserWarning: The `size` p
warnings.warn(msg, UserWarning)

```
Out[28]: <seaborn.axisgrid.PairGrid at 0x1c85e039588>
```




Imputing Values

```
In [35]: #There is no need of date
data = data.drop(['date'], axis=1)
```

```
In [36]: #Checking there is any null value or not
plt.figure(figsize=(10, 5))
sns.heatmap(data.isnull())
```

```
Out[36]: <matplotlib.axes._subplots.AxesSubplot at 0x1c86ddf10f0>
```



There are no missing values in the data

1.4 Model Training

This is a Regression problem since we are predicting house prices which is a continuous random variable. The steps involved are as follows:

1. Standardize or Normalize Training Data
2. Train Test Split
3. Train Model
4. Evaluation based on RMSE

Note: We will be implementing two models, Linear Regression and Random Forest Regressor

```
In [37]: # Take target variable into y
         y = data['price']
```

```
In [38]: # Delete the price
         del data['price']
```

```
In [39]: #Take their values in X and y
         X = data.values
         y = y.values
```

```
In [40]: # Split data into train and test formate
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=
```

Evaluation Metric

```
In [43]: def get_rmse(y_pred, y_target):
        return np.sqrt(np.mean(np.square(y_pred.reshape(-1,) - y_target.reshape(-1,))))
```

Linear Regression

```
In [41]: #Train the model
from sklearn import linear_model
model = linear_model.LinearRegression()
```

```
In [42]: #Fit the model
model.fit(X_train, y_train)
```

```
Out[42]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
In [44]: y_pred = model.predict(X_test)
```

```
In [46]: # Score/Accuracy
print("Accuracy --> ", model.score(X_test, y_test)*100)
```

```
Accuracy --> 75.90091673275158
```

```
In [45]: get_rmse(y_pred, y_test)
```

```
Out[45]: 0.25540977262992337
```

Random Forest Regression

```
In [47]: #Train the model
from sklearn.ensemble import RandomForestRegressor
model = RandomForestRegressor(n_estimators=1000)
```

```
In [48]: #Fit the model
model.fit(X_train, y_train)
```

```
Out[48]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
                                max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=1000,
                                n_jobs=None, oob_score=False, random_state=None,
                                verbose=0, warm_start=False)
```

```
In [49]: #Score/Accuracy
print("Accuracy --> ", model.score(X_test, y_test)*100)
```


Accuracy --> 88.56592044444415

```
In [50]: y_pred = model.predict(X_test)
```

```
In [51]: get_rmse(y_pred, y_test)
```

```
Out[51]: 0.17592917014179624
```